

GRADIENT BOOSTING ALGORITHM

This code uses the **Gradient Boosting algorithm** to predict the `PRODUCT_LENGTH` of a product based on its `PRODUCT_TYPE_ID`. Here is a step-by-step explanation of the code:

- The necessary libraries are imported: `pandas` to work with dataframes, `GradientBoostingRegressor` from `sklearn.ensemble` to build the gradient boosting model, and `mean_absolute_percentage_error` from `sklearn.metrics` to evaluate the model's performance.
- The train and test data are loaded using `pandas read_csv()` function. The data is stored in `pandas dataframes train_df` and `test_df`, respectively.
- The train and test data are concatenated using `pd.concat()` function to form a `full_df` dataframe. The `NaN` values are replaced with 0 in the `full_df` dataframe using the `fillna()` method.
- The features are selected from the train and test data, and `X_train`, `X_test`, `y_train`, and `y_test` are created as separate dataframes using `pandas`.
- A `GradientBoostingRegressor()` object called `gb` is created, and the `fit()` method is used to train the model using the `X_train` and `y_train` data. The model is used to make predictions on the test data using the `predict()` method. The predicted values are stored in `y_pred`.
- The `mean_absolute_percentage_error()` function is used to calculate the model's accuracy by comparing the actual values (`y_test`) with the predicted values (`y_pred`). The accuracy is stored in the `score` variable.
- The predicted values are saved in a new `pandas dataframe` called `sub_df`, along with their corresponding `PRODUCT_IDs`. The index of the `sub_df` dataframe is set to `PRODUCT_ID` using the `set_index()` method. The data in `sub_df` is written to a new CSV file called 'Gradient Boosting Result Submission.csv' using the `to_csv()` method.

In terms of feature engineering, the code only uses the `PRODUCT_TYPE_ID` feature to predict the `PRODUCT_LENGTH`. No additional feature engineering is performed.